

## Walmart

Walmart is an American multinational retail corporation that operates a chain of supercenters, discount departmental stores, and grocery stores from the United States. Walmart has more than 100 million customers worldwide.

## Business Problem

The Management team at Walmart Inc. wants to analyze the customer purchase behavior (specifically, purchase amount) against the customer's gender and the various other factors to help the business make better decisions. They want to understand if the spending habits differ between male and female customers: Do women spend more on Black Friday than men? (Assume 50 million customers are male and 50 million are female).

## Dataset

The company collected the transactional data of customers who purchased products from the Walmart Stores during Black Friday. The dataset has the following features: Dataset link: Walmart\_data.csv

- User\_ID: User ID
- Product\_ID: Product ID
- Gender: Sex of User
- Age: Age in bins
- Occupation: Occupation(Masked)
- City\_Category: Category of the City (A,B,C)
- StayInCurrentCityYears: Number of years stay in current city
- Marital\_Status: Marital Status
- ProductCategory: Product Category (Masked)
- Purchase: Purchase Amount

In [36]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

In [4]:

```
df = pd.read_csv('walmart_data.txt')
```

In [5]:

df

Out[5]:

|        | User_ID | Product_ID | Gender | Age   | Occupation | City_Category | Stay_In_Current_City_Ye |
|--------|---------|------------|--------|-------|------------|---------------|-------------------------|
| 0      | 1000001 | P00069042  | F      | 0-17  | 10         | A             |                         |
| 1      | 1000001 | P00248942  | F      | 0-17  | 10         | A             |                         |
| 2      | 1000001 | P00087842  | F      | 0-17  | 10         | A             |                         |
| 3      | 1000001 | P00085442  | F      | 0-17  | 10         | A             |                         |
| 4      | 1000002 | P00285442  | M      | 55+   | 16         | C             |                         |
| ...    | ...     | ...        | ...    | ...   | ...        | ...           |                         |
| 550063 | 1006033 | P00372445  | M      | 51-55 | 13         | B             |                         |
| 550064 | 1006035 | P00375436  | F      | 26-35 | 1          | C             |                         |
| 550065 | 1006036 | P00375436  | F      | 26-35 | 15         | B             |                         |
| 550066 | 1006038 | P00375436  | F      | 55+   | 1          | C             |                         |
| 550067 | 1006039 | P00371644  | F      | 46-50 | 0          | B             |                         |

550068 rows × 10 columns

In [14]:

# Checking structure &amp; D-types of our columns in data

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null int64
1   Product_ID                           550068 non-null object
2   Gender                                550068 non-null object
3   Age                                   550068 non-null object
4   Occupation                            550068 non-null int64
5   City_Category                         550068 non-null object
6   Stay_In_Current_City_Years           550068 non-null object
7   Marital_Status                        550068 non-null int64
8   Product_Category                     550068 non-null int64
9   Purchase                             550068 non-null int64
dtypes: int64(5), object(5)
memory usage: 42.0+ MB
```

In [15]:

```
df.shape
```

Out[15]:

```
(550068, 10)
```

## Our data contains 550k records & 10 columns

In [13]:

```
# Null-values check point
```

```
df.isnull().sum()
```

Out[13]:

```
User_ID          0
Product_ID       0
Gender           0
Age              0
Occupation       0
City_Category    0
Stay_In_Current_City_Years  0
Marital_Status   0
Product_Category 0
Purchase         0
dtype: int64
```

## There are no null values in our data

In [10]:

```
df.describe(include='O')
```

Out[10]:

|        | Product_ID | Gender | Age    | City_Category | Stay_In_Current_City_Years |
|--------|------------|--------|--------|---------------|----------------------------|
| count  | 550068     | 550068 | 550068 | 550068        | 550068                     |
| unique | 3631       | 2      | 7      | 3             | 5                          |
| top    | P00265242  | M      | 26-35  | B             | 1                          |
| freq   | 1880       | 414259 | 219587 | 231173        | 193821                     |

In [11]:

```
df.describe()
```

Out[11]:

|       | User_ID      | Occupation    | Marital_Status | Product_Category | Purchase      |
|-------|--------------|---------------|----------------|------------------|---------------|
| count | 5.500680e+05 | 550068.000000 | 550068.000000  | 550068.000000    | 550068.000000 |
| mean  | 1.003029e+06 | 8.076707      | 0.409653       | 5.404270         | 9263.968713   |
| std   | 1.727592e+03 | 6.522660      | 0.491770       | 3.936211         | 5023.065394   |
| min   | 1.000001e+06 | 0.000000      | 0.000000       | 1.000000         | 12.000000     |
| 25%   | 1.001516e+06 | 2.000000      | 0.000000       | 1.000000         | 5823.000000   |
| 50%   | 1.003077e+06 | 7.000000      | 0.000000       | 5.000000         | 8047.000000   |
| 75%   | 1.004478e+06 | 14.000000     | 1.000000       | 8.000000         | 12054.000000  |
| max   | 1.006040e+06 | 20.000000     | 1.000000       | 20.000000        | 23961.000000  |

In [18]:

```
print(f'Number of Users in the data : {df.User_ID.nunique()}')
print(f'Number of Products in the data : {df.Product_ID.nunique()}')
```

Number of Users in the data : 5891

Number of Products in the data : 3631

## Distribution & Unique values in categorical columns of our data

In [32]:

```
df['Gender'].value_counts()/len(df)
```

Out[32]:

M 0.753105

F 0.246895

Name: Gender, dtype: float64

***Our data has 75% Male & 25% Female***

In [34]:

```
df['Age'].value_counts()/len(df)*100
```

Out[34]:

```
26-35    39.919974
36-45    19.999891
18-25    18.117760
46-50     8.308246
51-55     6.999316
55+       3.909335
0-17      2.745479
Name: Age, dtype: float64
```

**Majority of our data are from 26-35 Age.**

**Approx. 80% our data are >18 age.**

In [30]:

```
round(df['City_Category'].value_counts()/len(df)*100, 2)
```

Out[30]:

```
B    42.03
C    31.12
A    26.85
Name: City_Category, dtype: float64
```

In [31]:

```
round(df['Stay_In_Current_City_Years'].value_counts()/len(df)*100, 2)
```

Out[31]:


```
1    35.24
2    18.51
3    17.32
4+   15.40
0    13.53
Name: Stay_In_Current_City_Years, dtype: float64
```

## Uni-variate Analysis

**Understanding the distribution of data and detecting abnormal data points for continuous variables¶**


In [37]:

```
plt.figure(figsize=(10, 6))  
sns.histplot(data=df, x='Purchase', kde=True)  
plt.show()
```



In [38]:

```
sns.boxplot(data=df, x='Purchase', orient='h')  
plt.show()
```



**As we can see here from the plot Purchase column is having outliers above 20k**


## Understanding Distribution for Categorical columns

In [60]:

```
sns.countplot(data=df, x='Gender')
```

Out[60]:

```
<AxesSubplot:xlabel='Gender', ylabel='count'>
```




Our data has most of the users as Male.

In [61]:

```
sns.countplot(data=df, x='Occupation')
```

Out[61]:

```
<AxesSubplot:xlabel='Occupation', ylabel='count'>
```




There are 20 different types of Occupation and Product\_Category.

In [62]:

```
sns.countplot(data=df, x='City_Category')
```

Out[62]:

<AxesSubplot:xlabel='City\_Category', ylabel='count'>



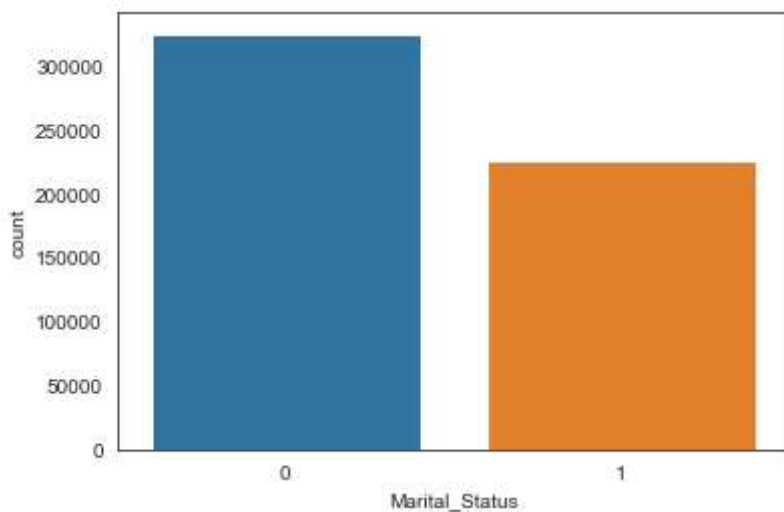
**More users belong to B City\_Category**

In [63]:

```
sns.countplot(data=df, x='Marital_Status')
```

Out[63]:

<AxesSubplot:xlabel='Marital\_Status', ylabel='count'>



**More users are Single as compared to Married**




In [64]:

```
sns.countplot(data=df, x='Product_Category')
```

Out[64]:

```
<AxesSubplot:xlabel='Product_Category', ylabel='count'>
```



**Product\_Category - 1, 5, 8, & 11 have highest purchasing frequency.**


In [40]:

```
fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(12, 8))

data = df['Age'].value_counts(normalize=True)*100
palette_color = sns.color_palette('BrBG_r')
axs[0].pie(x=data.values, labels=data.index, autopct='%0f%%', colors=palette_color)
axs[0].set_title("Age")

data = df['Stay_In_Current_City_Years'].value_counts(normalize=True)*100
palette_color = sns.color_palette('YlOrRd_r')
axs[1].pie(x=data.values, labels=data.index, autopct='%0f%%', colors=palette_color)
axs[1].set_title("Stay_In_Current_City_Years")

plt.show()
```




## Bi-variate Analysis

In [75]:

```
def box_plot_(col):
    sns.set_style("white")
    sns.boxplot(data=df, y='Purchase', x=col, palette='Set3').set_title(f"Purchase vs {col}")
```


In [76]:

```
box_plot_('Gender')
```




In [77]:

```
box_plot_('Age')
```



In [78]:

```
box_plot_('Occupation')
```




In [79]:

```
box_plot_('City_Category')
```



In [80]:

```
box_plot_('Product_Category')
```



In [81]:

```
box_plot_('Stay_In_Current_City_Years')
```



In [82]:

```
box_plot_('Marital_Status')
```



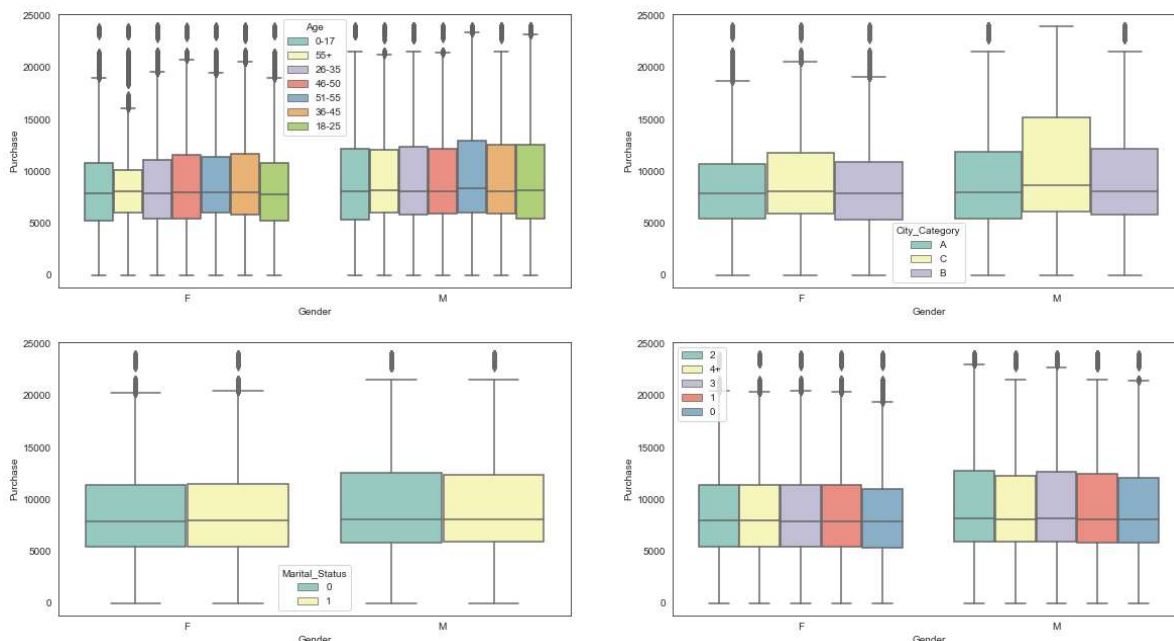
## Multivariate Analysis

In [44]:

```
fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(20, 6))
fig.subplots_adjust(top=1.5)
sns.boxplot(data=df, y='Purchase', x='Gender', hue='Age', palette='Set3', ax=axs[0,0])
sns.boxplot(data=df, y='Purchase', x='Gender', hue='City_Category', palette='Set3', ax=axs[0,1])

sns.boxplot(data=df, y='Purchase', x='Gender', hue='Marital_Status', palette='Set3', ax=axs[1,0])
sns.boxplot(data=df, y='Purchase', x='Gender', hue='Stay_In_Current_City_Years', palette='Set3', ax=axs[1,1])
axs[1,1].legend(loc='upper left')

plt.show()
```



## Average amount spend per customer for Male and Female

In [47]:

```
pur_df = df.groupby(['User_ID', 'Gender'])[['Purchase']].sum()  
pur_df = pur_df.reset_index()  
pur_df
```

Out[47]:

|      | User_ID | Gender | Purchase |
|------|---------|--------|----------|
| 0    | 1000001 | F      | 334093   |
| 1    | 1000002 | M      | 810472   |
| 2    | 1000003 | M      | 341635   |
| 3    | 1000004 | M      | 206468   |
| 4    | 1000005 | M      | 821001   |
| ...  | ...     | ...    | ...      |
| 5886 | 1006036 | F      | 4116058  |
| 5887 | 1006037 | F      | 1119538  |
| 5888 | 1006038 | F      | 90034    |
| 5889 | 1006039 | F      | 590319   |
| 5890 | 1006040 | M      | 1653299  |

5891 rows × 3 columns

In [48]:

```
# Gender wise value counts in pur_df  
pur_df['Gender'].value_counts()
```



Out[48]:

```
M    4225  
F    1666  
Name: Gender, dtype: int64
```

In [50]:

```
# histogram of average amount spend for each customer - Male & Female
pur_df[pur_df['Gender']=='M']['Purchase'].hist(bins=35)
plt.show()

pur_df[pur_df['Gender']=='F']['Purchase'].hist(bins=35)
plt.show()
```



In [51]:

```
male_avg = pur_df[pur_df['Gender']=='M']['Purchase'].mean()
female_avg = pur_df[pur_df['Gender']=='F']['Purchase'].mean()

print("Average amount spend by Male customers: {:.2f}".format(male_avg))
print("Average amount spend by Female customers: {:.2f}".format(female_avg))
```

Average amount spend by Male customers: 925344.40

Average amount spend by Female customers: 712024.39

**Male customers spend more money than female customers**



In [53]:

```

male_df = pur_df[pur_df['Gender']=='M']
female_df = pur_df[pur_df['Gender']=='F']

genders = ["M", "F"]

male_sample_size = 3000
female_sample_size = 1500
num_repitions = 1000
male_means = []
female_means = []

for _ in range(num_repitions):
    male_mean = male_df.sample(male_sample_size, replace=True)['Purchase'].mean()
    female_mean = female_df.sample(female_sample_size, replace=True)['Purchase'].mean()

    male_means.append(male_mean)
    female_means.append(female_mean)

```

In [54]:


```

fig, axis = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))

axis[0].hist(male_means, bins=35)
axis[1].hist(female_means, bins=35)
axis[0].set_title("Male - Distribution of means, Sample size: 3000")
axis[1].set_title("Female - Distribution of means, Sample size: 1500")

plt.show()

```



In [55]:

```

print("Population mean - Mean of sample means of amount spend for Male: {:.2f}".format(np.mean(male_means)))
print("Population mean - Mean of sample means of amount spend for Female: {:.2f}".format(np.mean(female_means)))

print("\nMale - Sample mean: {:.2f} Sample std: {:.2f}".format(male_df['Purchase'].mean(), male_df['Purchase'].std()))
print("Female - Sample mean: {:.2f} Sample std: {:.2f}".format(female_df['Purchase'].mean(), female_df['Purchase'].std()))

```

Population mean - Mean of sample means of amount spend for Male: 925262.26  
 Population mean - Mean of sample means of amount spend for Female: 713078.93

Male - Sample mean: 925344.40 Sample std: 985830.10  
 Female - Sample mean: 712024.39 Sample std: 807370.73

**Now using the Central Limit Theorem for the population we can say that:**

- Average amount spend by male customers is 925262.26
- Average amount spend by female customers is 713078.93

In [56]:

```
male_margin_of_error_clt = 1.96*male_df['Purchase'].std()/np.sqrt(len(male_df))
male_sample_mean = male_df['Purchase'].mean()
male_lower_lim = male_sample_mean - male_margin_of_error_clt
male_upper_lim = male_sample_mean + male_margin_of_error_clt

female_margin_of_error_clt = 1.96*female_df['Purchase'].std()/np.sqrt(len(female_df))
female_sample_mean = female_df['Purchase'].mean()
female_lower_lim = female_sample_mean - female_margin_of_error_clt
female_upper_lim = female_sample_mean + female_margin_of_error_clt

print("Male confidence interval of means: ({:.2f}, {:.2f})".format(male_lower_lim, male_upper_lim))
print("Female confidence interval of means: ({:.2f}, {:.2f})".format(female_lower_lim, female_upper_lim))
```

Male confidence interval of means: (895617.83, 955070.97)

Female confidence interval of means: (673254.77, 750794.02)

## Insights from married vs unmarried

In [57]:

```
amt_df = df.groupby(['User_ID', 'Marital_Status'])[['Purchase']].sum()
amt_df = amt_df.reset_index()

amt_df['Marital_Status'].value_counts()
```

Out[57]:

```
0    3417
1    2474
Name: Marital_Status, dtype: int64
```

In [58]:

```

marid_samp_size = 3000
unmarid_sample_size = 2000
num_repitions = 1000
marid_means = []
unmarid_means = []

for _ in range(num_repitions):
    marid_mean = amt_df[amt_df['Marital_Status']==1].sample(marid_samp_size, replace=True)[
    unmarid_mean = amt_df[amt_df['Marital_Status']==0].sample(unmarid_sample_size, replace=

    marid_means.append(marid_mean)
    unmarid_means.append(unmarid_mean)

fig, axis = plt.subplots(nrows=1, ncols=2, figsize=(20, 6))

axis[0].hist(marid_means, bins=35)

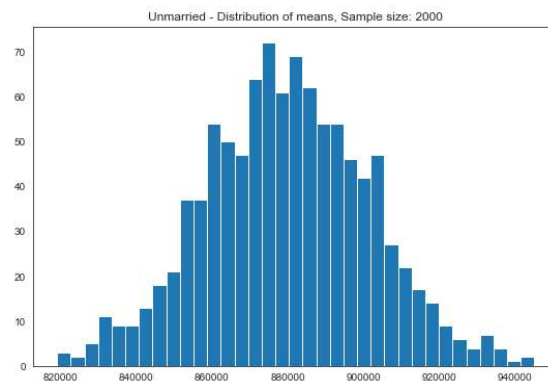
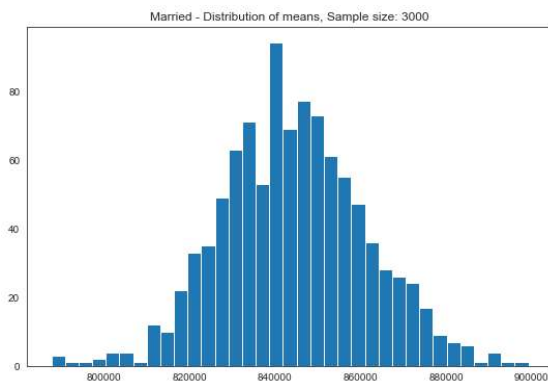
axis[1].hist(unmarid_means, bins=35)
axis[0].set_title("Married - Distribution of means, Sample size: 3000")
axis[1].set_title("Unmarried - Distribution of means, Sample size: 2000")

plt.show()

print("Population mean - Mean of sample means of amount spend for Married: {:.2f}".format(n
print("Population mean - Mean of sample means of amount spend for Unmarried: {:.2f}".format

print("\nMarried - Sample mean: {:.2f} Sample std: {:.2f}".format(
    amt_df[amt_df['Marital_Status']==1]['Purchase'].mean(), amt_df[amt_df['Marital_Status']]
print("Unmarried - Sample mean: {:.2f} Sample std: {:.2f}".format(
    amt_df[amt_df['Marital_Status']==0]['Purchase'].mean(), amt_df[amt_df['Marital_Status']]

```



Population mean - Mean of sample means of amount spend for Married: 844300.6

2

Population mean - Mean of sample means of amount spend for Unmarried: 87976

2.59

Married - Sample mean: 843526.80 Sample std: 935352.12

Unmarried - Sample mean: 880575.78 Sample std: 949436.25

In [59]:

```

for val in ["Married", "Unmarried"]:

    new_val = 1 if val == "Married" else 0

    new_df = amt_df[amt_df['Marital_Status']==new_val]

    margin_of_error_clt = 1.96*new_df['Purchase'].std()/np.sqrt(len(new_df))
    sample_mean = new_df['Purchase'].mean()
    lower_lim = sample_mean - margin_of_error_clt
    upper_lim = sample_mean + margin_of_error_clt

    print("{} confidence interval of means: {:.2f}, {:.2f}".format(val, lower_lim, upper_

```

Married confidence interval of means: (806668.83, 880384.76)

Unmarried confidence interval of means: (848741.18, 912410.38)

## Calculating the average amount spent by Age

In [85]:

```

df_age = df.groupby(['User_ID', 'Age'])[['Purchase']].sum()
df_age = df_age.reset_index()

df_age['Age'].value_counts()

```

Out[85]:

```

26-35    2053
36-45    1167
18-25    1069
46-50     531
51-55     481
55+       372
0-17      218
Name: Age, dtype: int64

```

In [90]:

```

sample_size = 200
num_repitions = 1000

all_means = {}

age_intervals = df_age['Age'].unique()
for age_interval in age_intervals:
    all_means[age_interval] = []

for age_interval in age_intervals:
    for _ in range(num_repitions):
        mean = df_age[df_age['Age']==age_interval].sample(sample_size, replace=True)['Purchase'].mean()
        all_means[age_interval].append(mean)

```

In [91]:

```

for val in df_age['Age'].unique():

    new_df = amt_df[amt_df['Age']==val]

    margin_of_error_clt = 1.96*new_df['Purchase'].std()/np.sqrt(len(new_df))
    sample_mean = new_df['Purchase'].mean()
    lower_lim = sample_mean - margin_of_error_clt
    upper_lim = sample_mean + margin_of_error_clt

    print("For age {} --> confidence interval of means: ({:.2f}, {:.2f})".format(val, lower

```

```

For age 0-17 --> confidence interval of means: (527662.46, 710073.17)
For age 55+ --> confidence interval of means: (476948.26, 602446.23)
For age 26-35 --> confidence interval of means: (945034.42, 1034284.21)
For age 46-50 --> confidence interval of means: (713505.63, 871591.93)
For age 51-55 --> confidence interval of means: (692392.43, 834009.42)
For age 36-45 --> confidence interval of means: (823347.80, 935983.62)
For age 18-25 --> confidence interval of means: (801632.78, 908093.46)

```

## Recommendations

- Men spent more money than women, So company should focus on retaining the male customers and getting more male customers.
- Product\_Category - 1, 5, 8, & 11 have highest purchasing frequency. it means these are the products in these categories are liked more by customers. Company can focus on selling more of these products or selling more of the products which are purchased less.
- Unmarried customers spend more money than married customers, So company should focus on acquisition of Unmarried customers.
- Customers in the age 18-45 spend more money than the others, So company should focus on acquisition of customers who are in the age 18-45
- Male customers living in City\_Category C spend more money than other male customers living in B or C, Selling more products in the City\_Category C will help the company increase the revenue.

In [ ]: